

Homework 4

Patrick Addona

February 15, 2025

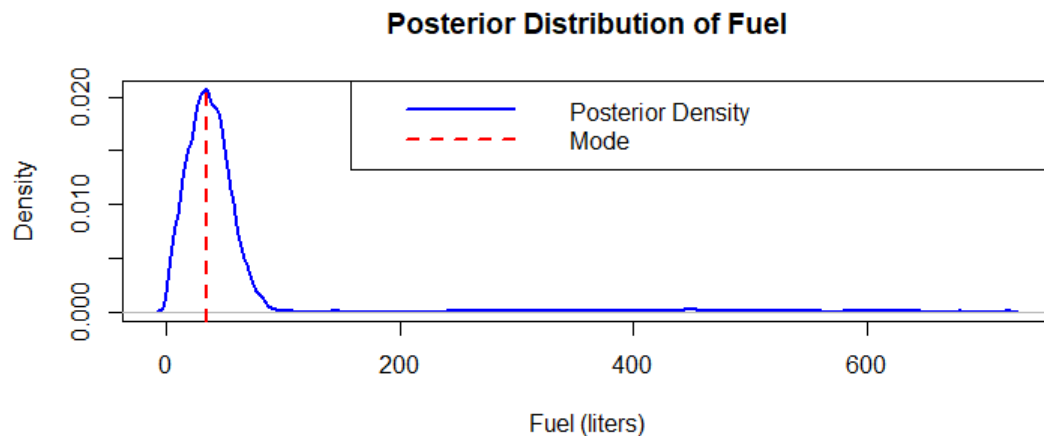
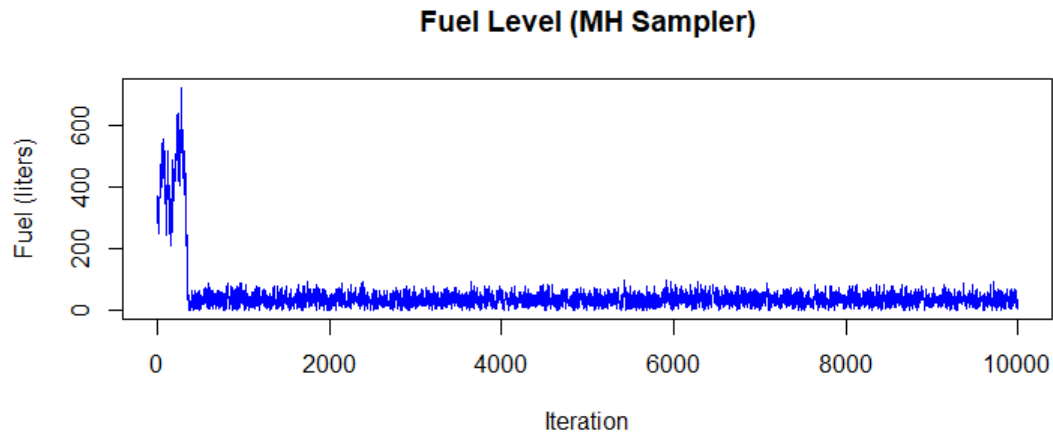
Question 1

Revisit the estimation problem about the fuel level (just this part) from the previous problem set using implementation of the Metropolis Hastings algorithm.

To this end, please: assess whether your numerical inference about the mode converges.

To implement the Metropolis Hastings algorithm in R, I found an introduction to the algorithm published by Ken Wood to be very helpful with the mathematical formulation. I chose to run the chain for 10,000 iterations (which we will see is more than enough), with a standard deviation of 30 to properly probe the posterior space. If I were to choose more iterations, the results would be redundant, and too few iterations may not guarantee convergence. Likewise, a standard deviation that is too high may never accept proposals, and one that is too low may never reject and take much longer to converge.

The results of the algorithm are shown in the following graphs:



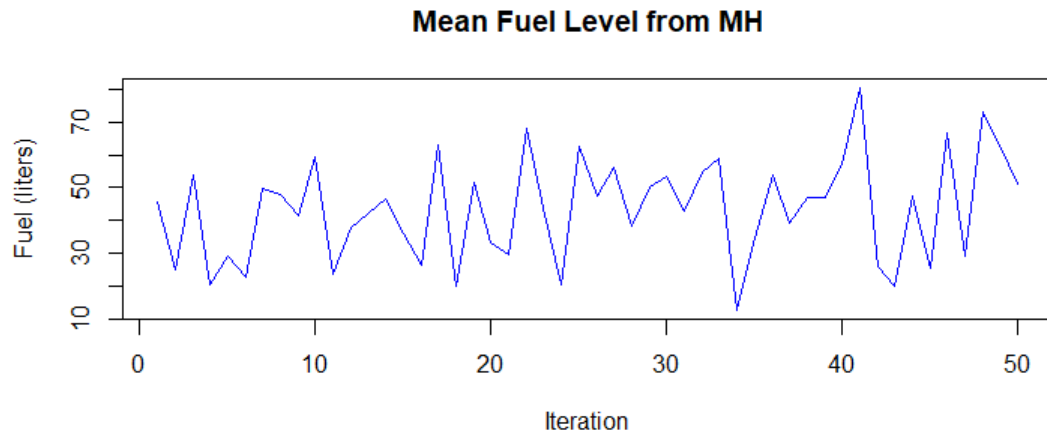
We see that after an initial burn-in period, the MH algorithm converges to the expected value of approximately 34, which is our observed fuel level. Likewise, the posterior distribution created has its peak at the mode of around 34, with the long right tail corresponding to the burn-in period. Were we to truncate this, it would appear cleaner. This analysis is reproducible through the use of a chosen seed.

Ken Wood's Metropolis-Hastings Algorithm page: <https://rpubs.com/ROARMarketingConcepts/1063733>

Question 2

Please define and use a positive control to assess the accuracy of your implemented method for the inference in Question 1.

To define a positive control, I chose to assume that we know the true fuel level is 40 liters. After generating a noisy observation "y_sim", I ran the MH algorithm 50 times and averaged the posterior means. The means themselves may be quite different, but the mean mean approaches the positive control value. This indicates that our method does indeed converge as expected, though noise is certainly present. Below the posterior means are shown:



As before, this analysis is reproducible with the use of designated seeds.

Question 3

Please compare the required number of runs needed for converged inferences in part b from the Metropolis Hastings algorithm with the Bayes Monte Carlo algorithm.

Because the Bayes Monte Carlo algorithm samples directly from the prior, we do not have to consider burn-in. We do not have to worry about a Markov Chain walking in the correct direction, which means the method is effective for low-dimensional problems. However, the sample size must be large enough so that an appropriate portion of them end up in the more likely regions. In my implementation from last assignment, I utilized 1,000,000 samples, though many less may have done the job.

In comparison, the MH algorithm needs to run for long enough to even accurately sample the posterior. I used 10,000 iterations, which ended up being plenty to account for burn-in. In general, MH would be more suited for a high-dimensional problem, as the samples required would not scale geometrically with dimension like in BMC.

Appendix

Problem 1:

```
1 ## author: Patrick Addona
2 ## copyright by the author
3 ## distributed under the GNU general public license
4 ## https://www.gnu.org/licenses/gpl.html
5 ## no warranty (see license details at the link above)
6
7 set.seed(1) # for reproducibility
8
9 # initial settings
10 n_iter <- 10000      # total iterations
11 sigma <- 30         # proposal standard deviation
12 y <- 34             # observed sensor reading
13 n_accept <- 0       # initialize to later calculate acceptance prob
14
15 # prior: Uniform(0, 182)
16 prior <- function(x) {
17   if (x >= 0 && x <= 182) {
18     return(1/182)
19   } else {
20     return(0)
21   }
22 }
23
24 # likelihood: N(y=34 | X, sd=20)
25 likelihood <- function(x) {
26   dnorm(y, mean = x, sd = 20, log = FALSE)
27 }
28
29 # initialize the chain
30 X <- numeric(n_iter)
31 X[1] <- 300 # choose value outside prior range to demonstrate convergence
32
33 # begin Metropolis-Hastings
34 for (t in 1:(n_iter-1)) {
35   # current state
36   x_curr <- X[t]
37
38   # propose x_star ~ N(x_curr, sigma^2)
39   x_star <- rnorm(1, mean = x_curr, sd = sigma)
40
41   # compute acceptance ratio
42   p_star <- prior(x_star) * likelihood(x_star)
43   p_curr <- prior(x_curr) * likelihood(x_curr)
44
45   # ratio, set alpha to 1 if p_curr = 0 to avoid division by 0 and encourage movement towards high-
46   # density regions
47   # alpha simplifies to p_star / p_curr because the likelihood is Gaussian
48   alpha <- min(1, p_star / p_curr)
49   if (p_curr == 0) {
50     alpha = 1
51   }
52
53   # accept/reject
54   if (runif(1) < alpha) {
55     X[t+1] <- x_star
56     n_accept <- n_accept + 1
57   } else {
58     X[t+1] <- x_curr
59   }
60 }
61
62 # remove burn-in
63 burn_in <- 0 #set to 0 to show convergence
64 X_post <- X[(burn_in+1):n_iter]
65
66 accept_rate <- n_accept / (n_iter - 1)
67 cat("Acceptance rate:", accept_rate, "\n")
```

```

68 # check the mode convergence
69 dens <- density(X_post)
70 mode_est <- dens$x[which.max(dens$y)]
71 cat("Posterior Mode estimate:", mode_est, "\n")
72
73 # plot the chain and the posterior density
74 par(mfrow=c(2,1))
75
76 # trace plot of chain values
77 plot(X_post, type = "l", col="blue",
78      main="Fuel Level (MH Sampler)",
79      xlab="Iteration", ylab="Fuel (liters)")
80
81 # posterior density estimate (after burn-in)
82 plot(dens, main="Posterior Distribution of Fuel",
83      xlab="Fuel (liters)", col="blue", lwd=2)
84 abline(v=mode_est, col="red", lty=2, lwd=2)
85
86 legend("topright", legend=c("Posterior Density", "Mode"),
87      col=c("blue", "red"), lty=c(1,2), lwd=2)

```

Problem 2:

```

1 ## author: Patrick Addona
2 ## copyright by the author
3 ## distributed under the GNU general public license
4 ## https://www.gnu.org/licenses/gpl.html
5 ## no warranty (see license details at the link above)
6
7 set.seed(1) # for reproducibility
8
9 true_fuel <- 40
10 n_reps <- 50
11 # generate a single noisy sensor reading from N(50, 20^2)
12 y_obs <- rnorm(1, mean = true_fuel, sd = 20)
13 n_iter <- 10000
14 sigma <- 5
15 n_accept <- 0 # initialize to later calculate acceptance prob
16 posterior_means <- numeric(n_reps)
17
18 # prior: Uniform(0, 182)
19 prior <- function(x) {
20   if (x >= 0 && x <= 182) {
21     return(1/182)
22   } else {
23     return(0)
24   }
25 }
26
27 # likelihood: N(y=34 | X, sd=20)
28 likelihood <- function(x) {
29   dnorm(y_obs, mean = x, sd = 20, log = FALSE)
30 }
31
32 #repeat MH algorithm to show positive control is accurately reached, on average
33 for (rep in 1:n_reps){
34   y_obs <- rnorm(1, mean = true_fuel, sd = 20)
35   X <- numeric(n_iter)
36   X[1] <- 30
37
38   # begin Metropolis-Hastings
39   for (t in 1:(n_iter-1)) {
40     # current state
41     x_curr <- X[t]
42
43     # propose x_star ~ N(x_curr, sigma^2)
44     x_star <- rnorm(1, mean = x_curr, sd = sigma)
45
46     # compute acceptance ratio
47     p_star <- prior(x_star) * likelihood(x_star)
48     p_curr <- prior(x_curr) * likelihood(x_curr)
49

```

```

50 # ratio, set alpha to 1 if p_curr = 0 to avoid division by 0 and encourage movement towards
    high-density regions
51 # alpha simplifies to p_star / p_curr because the likelihood is Gaussian
52 alpha <- min(1, p_star / p_curr)
53 if (p_curr == 0) {
54   alpha = 1
55 }
56
57 # accept/reject
58 if (runif(1) < alpha) {
59   X[t+1] <- x_star
60   n_accept <- n_accept + 1
61 } else {
62   X[t+1] <- x_curr
63 }
64 }
65
66 # remove burn-in
67 burn_in <- 0 #set to 0 to show convergence
68 X_post <- X[(burn_in+1):n_iter]
69
70 accept_rate <- n_accept / (n_iter - 1)
71
72 posterior_means[rep] <- mean(X_post)
73 }
74
75 cat("Mean of posterior means =", mean(posterior_means), "\n")
76
77 # trace plot of chain values
78 plot(posterior_means, type = "l", col="blue",
79      main="Mean Fuel Level from MH",
80      xlab="Iteration", ylab="Fuel (liters)")

```